

Classifying influenza-related Tweets as influenza-aware or influenza influenza with Recurrent Neural Networks (Bidirectional GRU with Attention)

```
import os
import re
import json
import string
import logging
import multiprocessing

from utils import *
import numpy as np
import pandas as pd
from pylab import rcParams
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler

# Twitter preprocess, tokenizer
import twokenize
import preprocess_twitter
import preprocessor as p
p.set_options(p.OPT.URL, p.OPT.HASHTAG, p.OPT.EMOJI,
              p.OPT.SMILEY, p.OPT.RESERVED, p.OPT.MENTION)

def tokenize(row):
    punct = string.punctuation.replace("-", "")
    punct = string.punctuation.replace("'", "")
    exclude = set(string.punctuation)
    text = str(row)

    if text is not None:
        text = text.lower()
        text = ''.join(ch for ch in text if ch not in exclude)
        text = re.sub('[\'+punct+']', ' ', text)
        tokens = twokenize.tokenize(text)
        return tokens
    else:
        return None

# Keras utils
from keras.preprocessing.text import Tokenizer
from keras.utils.np_utils import to_categorical
from keras.preprocessing.sequence import pad_sequences

# Tensorflow
import tensorflow as tf
from tensorflow.contrib.rnn import GRUCell, LSTMCell
from tensorflow.python.ops.rnn import bidirectional_dynamic_rnn
from tensorflow.contrib.layers import fully_connected

# Convert labels to probability
cls2probs = lambda x: [1., 0] if x == 0 else [0., 1]

# Notebook related
from IPython.display import clear_output
from IPython.core.display import HTML
from IPython.display import display

logging.getLogger("tensorflow").setLevel(logging.WARNING)
pd.set_option('max_colwidth', 200)
%matplotlib inline
```

Loading tweets dataset

Awareness vs Infection Tweets Classification

- 0: Influenza infection
- 1: Influenza awareness

```
df = pd.read_csv("AwareInfection_all.csv", sep = "\t",
                 engine="c", header=None)
# Some oddly formatted lines
df = df[[0,1]]
df.dropna(inplace=True)
df[1] = df[1].astype(int)
# Shuffle df
df = df.sample(frac=1).reset_index()[[0,1]]
```

```
df.shape
```

```
(3151, 2)
```

```
df[1].value_counts()
```

```
1    1703
0    1448
Name: 1, dtype: int64
```

The Aware vs Infection dataset contains 3151 manually labelled data points. 1703 are labelled as 'Infection' while 1448 are labelled as 'Awareness'. This small difference does not imply an imbalanced class representation.

Pre-processing and Tokenizing Tweets

As recommended by [Haddi, 2013](#) we cleaned specific Twitter sequences, such as RT (retweet), user references (@user), emoticons (":", etc) and also URLs (http://*).

```
text = list(map(p.clean, df[0])) # Pre-processing
text = list(map(tokenize, text)) # Pre-processing
text = [" ".join(i) for i in text]

# Maximum vocabulary size
MAX_NB_WORDS = 6000

tokenizer = Tokenizer(num_words = MAX_NB_WORDS)
tokenizer.fit_on_texts(text)
sequences = tokenizer.texts_to_sequences(text)

word_index = tokenizer.word_index
print('%s unique tokens.' % len(word_index))
```

```
5559 unique tokens.
```

```
# Define end of sequence ("eos") as 0, to match zero-padding.
eos_id = 0

# Max sequence length (nr. of words)
MAX_SEQUENCE_LENGTH = 35

# zero-padding (post-sequence)
data = pad_sequences(sequences, padding="post",
                    maxlen=MAX_SEQUENCE_LENGTH)
```

Test-train split

We considered a 15%-85% test-train split.

```

labels = list(df[1])
labels = to_categorical(np.asarray(labels))
print('Data tensor shape:', data.shape)
print('Label tensor shape:', labels.shape)

VALIDATION_SPLIT = 0.15
np.random.seed(42)
indices = df.index.tolist()
np.random.shuffle(indices)
data = data[indices]
labels = labels[indices]
nb_validation_samples = int(VALIDATION_SPLIT * data.shape[0])

xx_train = data[:-nb_validation_samples]
yy_train = labels[:-nb_validation_samples]
xx_val = data[-nb_validation_samples:]
yy_val = labels[-nb_validation_samples:]

print('Training shape: ', xx_train.shape)
print('Validation shape: ', xx_val.shape)

```

```

Data tensor shape: (3151, 35)
Label tensor shape: (3151, 2)
Training shape: (2679, 35)
Validation shape: (472, 35)

```

Word embeddings

Word embeddings trained on Twitter data can be found in the [GloVe](#) project. These were trained on 2 billion tweets, comprising 27 billion tokens with a vocabulary length of around 1.2 million. We used the highest embedding dimension available, 200.

```

embeddings_index = {}

with open('glove.twitter.27B.200d.txt') as file:
    for line in file:
        values = line.split()
        word = values[0]
        coefs = np.asarray(values[1:], dtype='float32')
        embeddings_index[word] = coefs

print('Found %s word vectors.' % len(embeddings_index))

```

```

Found 1193514 word vectors.

```

```

# Word embeddings dimension
# GloVe Twitter 200d
EMBEDDING_DIM = 200

embedding_matrix = np.zeros((len(word_index) + 1, EMBEDDING_DIM))
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        # words not found in embedding index will be all-zeros.
        embedding_matrix[i] = embedding_vector
embeddings = embedding_matrix

```

Bidirectional RNN (GRU) with Attention Mechanism

Based on [Ilya's](#) TensorFlow implementation.

```

# Network Parameters
EMBED_DIM = 200
HIDDEN_SIZE = 32 # GRU
ATTENTION_SIZE = 256

```

```

NUM_CLASSES = 2

# Optimizer (Adam) parameters
LEARNING_RATE = 0.000175
EPSILON = 1e-4
BETA1 = 0.9
BETA2 = 0.99

# L2 regularization coefficient
BETA = 0

# Probability of keeping a neuron
DROPOUT = 0.5

# Batch size and epochs
BATCH_SIZE = 96
EPOCHS = 32

```

```

tf.reset_default_graph()

# Placeholders
batch_ph = tf.placeholder(tf.int32, [None, MAX_SEQUENCE_LENGTH])
target_ph = tf.placeholder(tf.float32, [None, NUM_CLASSES])
seq_len_ph = tf.placeholder(tf.int32, [None])
keep_prob_ph = tf.placeholder(tf.float32)
embeddings_ph = tf.placeholder(tf.float32, [len(word_index)+1, EMBED_DIM])

```

```

# Embedding layer
embeddings_ph = tf.placeholder(tf.float32, [len(word_index)+1, EMBED_DIM])
embeddings_var = tf.Variable(tf.constant(0., shape=[len(word_index)+1, EMBED_DIM]),
                             trainable=False)
init_embeddings = embeddings_var.assign(embeddings_ph)
batch_embedded = tf.nn.embedding_lookup(embeddings_var, batch_ph)

# Bi-RNN layer
outputs, _ = tf.nn.bidirectional_dynamic_rnn(GRUCell(HIDDEN_SIZE), GRUCell(HIDDEN_SIZE),
                                              inputs=batch_embedded, sequence_length=seq_len_ph,
                                              dtype=tf.float32, scope="bi_rnn1")

outputs = tf.concat(outputs, 2)

# Attention mechanism
W_omega = tf.Variable(tf.random_normal([2 * HIDDEN_SIZE, ATTENTION_SIZE], stddev=0.1))
b_omega = tf.Variable(tf.random_normal([ATTENTION_SIZE], stddev=0.1))
u_omega = tf.Variable(tf.random_normal([ATTENTION_SIZE], stddev=0.1))

v = tf.tanh(tf.matmul(tf.reshape(outputs, [-1, 2 * HIDDEN_SIZE]), W_omega) +
             tf.reshape(b_omega, [1, -1]))

vu = tf.matmul(v, tf.reshape(u_omega, [-1, 1]))
exps = tf.reshape(tf.exp(vu), [-1, MAX_SEQUENCE_LENGTH])
alphas = exps / tf.reshape(tf.reduce_sum(exps, 1), [-1, 1])

# Output of Bi-RNN reduced with attention vector
output = tf.reduce_sum(outputs * tf.reshape(alphas, [-1, MAX_SEQUENCE_LENGTH, 1]), 1)

# Dropout
drop = tf.nn.dropout(output, keep_prob_ph)

# Fully connected layer
W = tf.Variable(tf.truncated_normal([HIDDEN_SIZE * 2, NUM_CLASSES], stddev=0.1), name="W")
b = tf.Variable(tf.constant(0., shape=[NUM_CLASSES]), name="b")
y_hat = tf.nn.xw_plus_b(output, W, b, name="scores")

# Loss function
cross_entropy = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(labels=target_ph,
                                                                           logits=y_hat),
                               name="cross_entropy")
l2_loss = tf.nn.l2_loss(W, name="l2_loss")
loss = cross_entropy + l2_loss * BETA

# Optimizer
optimizer = tf.train.AdamOptimizer(learning_rate=LEARNING_RATE,
                                    beta1=BETA1, beta2=BETA2,
                                    epsilon=EPSILON).minimize(loss)

# Accuracy
accuracy = tf.reduce_mean(tf.cast(tf.equal(tf.argmax(target_ph, 1),
                                                    tf.argmax(y_hat, 1)),
                               tf.float32))

```

```

# Initialize mode
train_batch_generator = batch_generator(xx_train, yy_train, BATCH_SIZE)

loss_tr_l = []
loss_val_l = []
ce_tr_l = [] # Cross-entropy
ce_val_l = []
acc_tr_l = [] # Accuracy
acc_val_l = []

with tf.Session() as sess:

    sess.run(tf.global_variables_initializer())
    sess.run(init_embeddings, feed_dict={embeddings_ph: embeddings})

    for epoch in range(EPOCHS):
        for i in range(int(xx_train.shape[0] / BATCH_SIZE)):

            x_batch, y_batch = train_batch_generator.__next__()

            # Using variable sequence length
            seq_len_tr = np.array([list(x).index(eos_id) + 1 for x in x_batch])
            sess.run(optimizer, feed_dict={batch_ph: x_batch, target_ph: y_batch,
                                           seq_len_ph: seq_len_tr, keep_prob_ph: DROPOUT})

            y_pred_tr, ce_tr, loss_tr, acc_tr = sess.run([y_hat, cross_entropy, loss, accuracy],
                                                         feed_dict={batch_ph: x_batch,
                                                                       target_ph: y_batch,
                                                                       seq_len_ph: seq_len_tr,
                                                                       keep_prob_ph: 1.0})

            y_pred_val, ce_val, loss_val, acc_val = [], 0, 0, 0
            num_val_batches = int(xx_val.shape[0] / BATCH_SIZE)
            for i in range(num_val_batches):
                x_batch_val, y_batch_val = xx_val[i * BATCH_SIZE : (i + 1) * BATCH_SIZE], \
                                           yy_val[i * BATCH_SIZE : (i + 1) * BATCH_SIZE]
                seq_len_val = np.array([list(x).index(eos_id) + 1 for x in x_batch_val])
                y_pred_val_, ce_val_, loss_val_, acc_val_ = sess.run([y_hat,
                                                                       cross_entropy,
                                                                       loss, accuracy],
                                                                       feed_dict={batch_ph: x_batch_val,
                                                                       target_ph: y_batch_val,
                                                                       seq_len_ph: seq_len_val,
                                                                       keep_prob_ph: 1.0})

                y_pred_val += list(y_pred_val_)
                ce_val += ce_val_
                loss_val += loss_val_
                acc_val += acc_val_

            y_pred_val = np.array(y_pred_val)
            ce_val /= num_val_batches
            loss_val /= num_val_batches
            acc_val /= num_val_batches

            y_pred_tr = np.array([cls2probs(cls) for cls in np.argmax(y_pred_tr, 1) - 1])
            y_pred_val = np.array([cls2probs(cls) for cls in np.argmax(y_pred_val, 1) - 1])

            loss_tr_l.append(loss_tr)
            loss_val_l.append(loss_val)
            ce_tr_l.append(ce_tr)
            ce_val_l.append(ce_val)
            acc_tr_l.append(acc_tr)
            acc_val_l.append(acc_val)

        clear_output(wait=True)
        print("epoch: {}".format(epoch))
        print("\t Train loss: {:.3f}\t ce: {:.3f}\t acc: {:.3f}".format(
            loss_tr, ce_tr, acc_tr))
        print("\t Valid loss: {:.3f}\t ce: {:.3f}\t acc: {:.3f}".format(
            loss_val, ce_val, acc_val))

    plt.figure(figsize=(7,5))
    plt.plot(ce_tr_l, color='blue', label='ce_tr')
    plt.plot(ce_val_l, color='red', label='ce_val')
    plt.xlim(0, EPOCHS - 1)
    plt.ylim(0, 1)
    plt.legend()
    plt.show()

```

```

results = [acc_val]


saver = tf.train.Saver()
saver.save(sess, 'BiGRU_att_1')

```

```

epoch: 31
    Train loss: 0.334      ce: 0.334      acc: 0.854
    Valid loss: 0.320      ce: 0.320      acc: 0.859


```



```

plt.figure(figsize=(7,5))
plt.plot(acc_tr_l, color='blue', label='Training acc')
plt.plot(acc_val_l, color='red', label='Validation acc')
plt.xlim(0, EPOCHS - 1)
plt.ylim(0, 1)
plt.legend()
plt.show()

```



After 31 epochs we attained ~86% of validation accuracy (85% training accuracy), with similar cross-entropy levels (around 0.32-0.33) in both sets.

Classification of influenza-related tweets

The trained RNN is now used to classify a set of around 800 thousand English influenza-related tweets extracted from the [SNAP](#) Twitter dataset.

```

def return_preds(text):
    all_preds = []
    top_words = []

    # Pre-processing
    text = list(map(p.clean, text))
    text = list(map(tokenize, text))
    text = [" ".join(i) for i in text]
    input_seqs = tokenizer.texts_to_sequences(text)

```

```

input_seqs = pad_sequences(input_seqs, maxlen=MAX_SEQUENCE_LENGTH, padding='post')

# Get predictions and top words
with tf.Session() as sess:
    new_saver = tf.train.import_meta_graph('BiGRU_att_1.meta')
    new_saver.restore(sess, tf.train.latest_checkpoint('./'))
    for i in input_seqs:
        x_batch = [i]
        seq_len = np.array([MAX_SEQUENCE_LENGTH])
        y_pred, alphas_pred = sess.run([y_hat, alphas], feed_dict={batch_ph: x_batch,
                                                                    seq_len_ph: seq_len,
                                                                    keep_prob_ph: 1.0})

        all_preds.append(y_pred)
        top_words.append(i[np.argmax(alphas_pred)])

# Convert softmax output to class, convert word index to word
all_preds = [np.argmax(i) for i in all_preds]
top_words = [list(word_index.keys())[list(word_index.values()).index(i)]
              if i > 0 else 'None' for i in top_words]

return all_preds, top_words

```

```

# Load english flu-related tweets (from SNAP7 Twitter dataset)
topred = pd.read_csv("/home/user/Thesis/FluTweets2009-EN.csv", sep="\t", header=None)

```

```

%%time
preds, top = return_preds(topred[1].tolist())

CPU times: user 5h 52min 17s, sys: 1h 19min 41s, total: 7h 11min 58s
Wall time: 1h 34min 42s

```

```

topred['preds'] = preds
topred['topword'] = top

```

```

topred.preds.value_counts()

1    606163
0    212135
Name: preds, dtype: int64

```

~3/4 of the tweets were classified as awareness, ~1/4 classified as infection.

```

topred.index = pd.to_datetime(topred[0], infer_datetime_format=True)
topred.drop(0, axis=1, inplace=True)

```


We show the absolute count of classified tweets overtime.

```

ax = topred[topred['preds']==1]['05-2009':'01-2010'].groupby(pd.TimeGrouper('W')).size().plot()
topred[topred['preds']==0]['05-2009':'01-2010'].groupby(pd.TimeGrouper('W')).size().plot(ax = ax)

<matplotlib.axes._subplots.AxesSubplot at 0x7f041f39cb38>

```



```
prop = pd.DataFrame(topred[topred['preds']==1]['05-2009':'01-2010'].groupby(pd.TimeGrouper('W')).size())
prop.columns = ['Awareness']
prop = prop.join(pd.DataFrame(topred[topred['preds']==0]['05-2009':'01-2010'].groupby(pd.TimeGrouper('W')).size()))
prop.columns = ['Awareness', 'Infection']
```

```
weekly = pd.read_csv("/home/user/Thesis/dailyTweets.csv", sep="\t", header=None)
weekly.index = pd.to_datetime(weekly[0], infer_datetime_format=True)
weekly = weekly.groupby(pd.TimeGrouper('W')).sum()
```

```
prop = prop.join(weekly)
prop.index.name = "Month"
prop['Awareness %'] = prop['Awareness'] / prop[1] * 100
prop['Infection %'] = prop['Infection'] / prop[1] * 100
prop[['Awareness %', 'Infection %']].to_csv("/home/user/Thesis/AwInf_prop.csv", sep="\t")
```

```
monthly_prop = prop[['Awareness', 'Infection', 1]].groupby(pd.TimeGrouper('M')).sum()
monthly_prop['Awareness %'] = monthly_prop['Awareness'] / monthly_prop[1] * 100
monthly_prop['Infection %'] = monthly_prop['Infection'] / monthly_prop[1] * 100
monthly_prop.to_csv("/home/user/Thesis/Twitter_AwInf_monthly.csv", sep="\t")
```

Proportion of classified tweets over time.

```
ax = prop[['Awareness %', 'Infection %']].plot(figsize=(11,5))
ax.set_ylabel("% of total Tweets")
fig = ax.get_figure()
fig.savefig('/home/user/Thesis/Twitter_AWvsINF.pdf', format='pdf')
```



```
prop.index = prop.index.shift(freq='W', n=-1)
```


The top words (with the most weight) for classification can now be extracted.

```
topred.groupby("preds")['topword'].value_counts()[0].head(15).tail(-1)
```

```
topword
got      21341
sick     21321
have     13697
getting  6381
up       6143
feel     4370
feeling  4363
from     4348
been     3712
down     3562
had       3166
fever    2890
caught   2750
swine    2487
Name: topword, dtype: int64
```

```
topred.groupby("preds")['topword'].value_counts()[1].head(15).tail(-1)
```

```
topword
vaccine    71022
swine      42119
shot       35655
shots      16682
from       13814
deaths     9800
vaccination 8864
cases      8447
health     7587
influenza  5566
up         5441
have       4773
out        4043
vaccines   3965
Name: topword, dtype: int64
```

```
import scipy.stats as stats
```

```
topred[(topred['topword']!='flu') & (topred['preds']==0)][['topword']]['2009-06':'2009-12'].groupby(
    pd.TimeGrouper('M')).apply(
    lambda x: x.value_counts().head(2))
```

```
2009-06-30 got      284
           sick     270
2009-07-31 got     1061
           sick     724
2009-08-31 got     4462
           sick     3764
2009-09-30 sick     5441
           got      5328
2009-10-31 sick     6498
           got      5487
2009-11-30 sick     3270
           got      3222
2009-12-31 got      1497
           sick     1354
Name: topword, dtype: int64
```

```
topred[(topred['topword']!='flu') & (topred['preds']==1)][['topword']]['2009-06':'2009-12'].groupby(
    pd.TimeGrouper('M')).apply(
    lambda x: x.value_counts().head(2))
```

```

2009-06-30  swine      2425
              cases      1499
2009-07-31  vaccine    6051
              swine      5306
2009-08-31  swine      14115
              vaccine    12005
2009-09-30  vaccine    14371
              shot       8577
2009-10-31  vaccine    21296
              shot       14389
2009-11-30  vaccine    11324
              shot       7248
2009-12-31  vaccine    4911
              shot       2684
Name: topword, dtype: int64


```

Wordclouds of most most decisive words for each classification

```
aw_text = " ".join([i for i in topred[topred['preds']==1]['topword'].tolist() if i != "flu"])
inf_text = " ".join([i for i in topred[topred['preds']==0]['topword'].tolist() if i != "flu"])
```

```
from wordcloud import WordCloud
wordcloud_aw = WordCloud(random_state=42, min_font_size=9, width=1024, height=768,
                          background_color='white', margin=1,
                          colormap='Blues_r',
                          max_words=90,
                          normalize_plurals=False,
                          collocations=False).generate(aw_text)

plt.figure(figsize=(9,6))
ax = plt.imshow(wordcloud_aw, interpolation="spline16")
plt.axis("off")
plt.show()
fig = ax.get_figure()
fig.savefig('/home/user/Thesis/Awareness_WordCloud.pdf', format='pdf')
```



```
wordcloud_inf = WordCloud(random_state=42, margin=1, width=1024, height=768,
                           background_color='white', min_font_size=9,
                           colormap='Reds_r',
                           max_words=90,
                           normalize_plurals=False,
                           collocations=False).generate(inf_text)

plt.figure(figsize=(9,6))
ax = plt.imshow(wordcloud_inf, interpolation="spline16")
plt.axis("off")
plt.show()
fig = ax.get_figure()
fig.savefig('/home/user/Thesis/Infection_WordCloud.pdf', format='pdf')
```

